



Dive Deeper in Finance


GTC Inspired Deep Learning Event, Munich

Daniel Egloff
Dr. sc. math.
Managing Director QuantAlea
October 10, 2017

Reinforcement Learning


Reinforcement Learning

- Emerged from two streams of research
 - Bellman's 1957 work on optimal control
 - Animal learning with trial and error
- Deep Learning allows RL to scale
- Recent applications
 - Learning to play Atari games
 - Google AlphaGo



Deep RL is Hot Topic in AI

From 1603 in 2016 to 2647 publications in first 3 quarters of 2017

 Semantic Scholar

deep reinforcement learning

SIGN IN ?

About 10,900 results

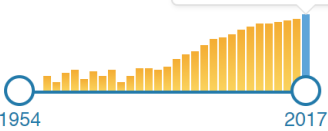
Sort by: Relevance ▾

Filter Results:

☐ Full text PDF available (9.602)

Publication Year

2.647 Publications – 2016 to 2017



1954 2017

This year (1.422)
Last 5 years (5.543)
Last 10 years (8.101)

Publication Type ▾

Author ▾


Journals and Conferences ▾

Deep Reinforcement Learning that Matters Trending





er Henderson, Riashat Islam, Philip Bachman, Joelle Pineau, Doina Precup, David Meger • 2017

Recent years, significant progress has been made in solving challenging problems across various domains using **deep reinforcement learning** (RL). Reproducing existing work and accurately judging the improvements offered by novel methods is vital to maintaining this rapid progress. Unfortunately, reproducing results for state-of-the-art **deep RL** methods is... [\(More\)](#)

Mentioned in 38 tweets

 **Vineet Vashishta**
@v_vashishta [Follow](#)

Deep Reinforcement Learning that Matters
arxiv.org/pdf/1709.06560... #DeepLearning #DRL
#machinelearning
12:00 AM - Sep 23, 2017

  22  20 


RL in Industrial Applications

- Smart robots for manufacturing, ...
- Inventory management
- Supply chain optimization
- Self driving cars
- Autonomous vehicles and drones
- Power grid management



RL in Finance?

- Optimize order execution
 - Learning to place limit orders to reach a position goal
 - Optimally distribute volume over time
 - Incorporate market impact in decision process





RL in Finance?

- Learning to trade
 - Risk and reward based trading strategy
 - Learned end-to-end from historical data
 - Online adaption to new market information
 - Consistently incorporate different information sources




More RL Applications in Finance

- Portfolio construction
- Manage customer lifetime value
- Transaction anomaly detection
- Funding and capital optimization
-



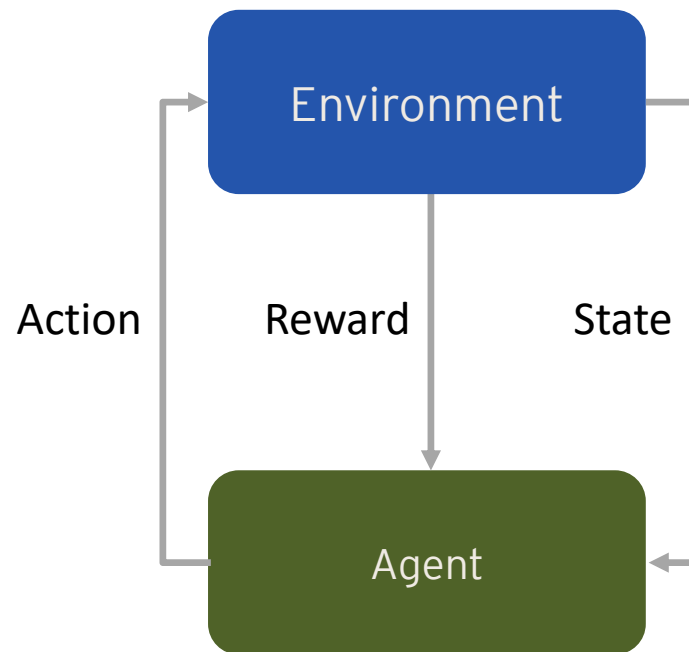
Customer Lifetime Value



RL Mechanics

RL Setup

- Learning a behavioral strategy which maximizes long term sum of rewards by a direct interaction with an unknown and uncertain environment



While not terminal do:

Agent perceives state s_t
Agent performs action a_t
Agent receives reward r_t
Environment evolves to state s_{t+1}

Markov Decision Process

- RL is a Markov Decision Process
 - Set of states \mathcal{S}
 - Set of actions \mathcal{A}
 - State transitions $p(s_{t+1}|s_t, a_t)$
 - Instantaneous reward function r_t
 - Policy $\pi : \mathcal{S} \rightarrow \mathcal{P}(\mathcal{A})$ maps states to probabilities over actions
 - Trajectory

$$s_0, a_0, r_0, s_1, a_1, r_1, s_2, \dots, a_t, r_t, s_{t+1}, \dots$$

$$a_t \sim \pi(s_t)$$

$$s_{t+1} \sim p(\cdot | s_t, a_t)$$

Markov Decision Process

- Episodic MDP
 - State resets after T steps
 - Cumulative discounted reward $R = \sum_{t=0}^{T-1} \gamma^t r_t$
 - Optimal policy

$$\pi^* = \operatorname{argmax}_{\pi} \mathbb{E}[R \mid \pi]$$

- Non-episodic MDP
- Average return

Markov Decision Process

- State value function
 - Expected return when starting in a state and following a policy

$$V^{\pi}(s) = \mathbb{E}[R \mid s, \pi]$$

- State action function
 - Expected return when starting in a state, taking an action and thereafter following a policy

$$Q^{\pi}(s, a) = \mathbb{E}[R \mid s, a, \pi]$$

$$V^{\pi}(s) = \max_a Q^{\pi}(s, a)$$

Markov Decision Process

- Dynamic programming
 - Bellman's equation

$$Q^{\pi}(s_t, a_t) = \mathbb{E}_{s_{t+1}}[r_t + \gamma Q^{\pi}(s_{t+1}, \pi(s_{t+1}))]$$

- Optimal Bellman equation

$$Q^*(s_t, a_t) = \mathbb{E}_{s_{t+1}} \left[r_t + \gamma \max_{a'} Q^*(s_{t+1}, a') \right]$$

RL Challenges

- Optimal policy must be inferred by trial and error interaction with environment
- Next observation depends on action of agent
- Observations may show strong temporal correlations
- Action often manifests itself after many state transitions
- Rewards may be sparse

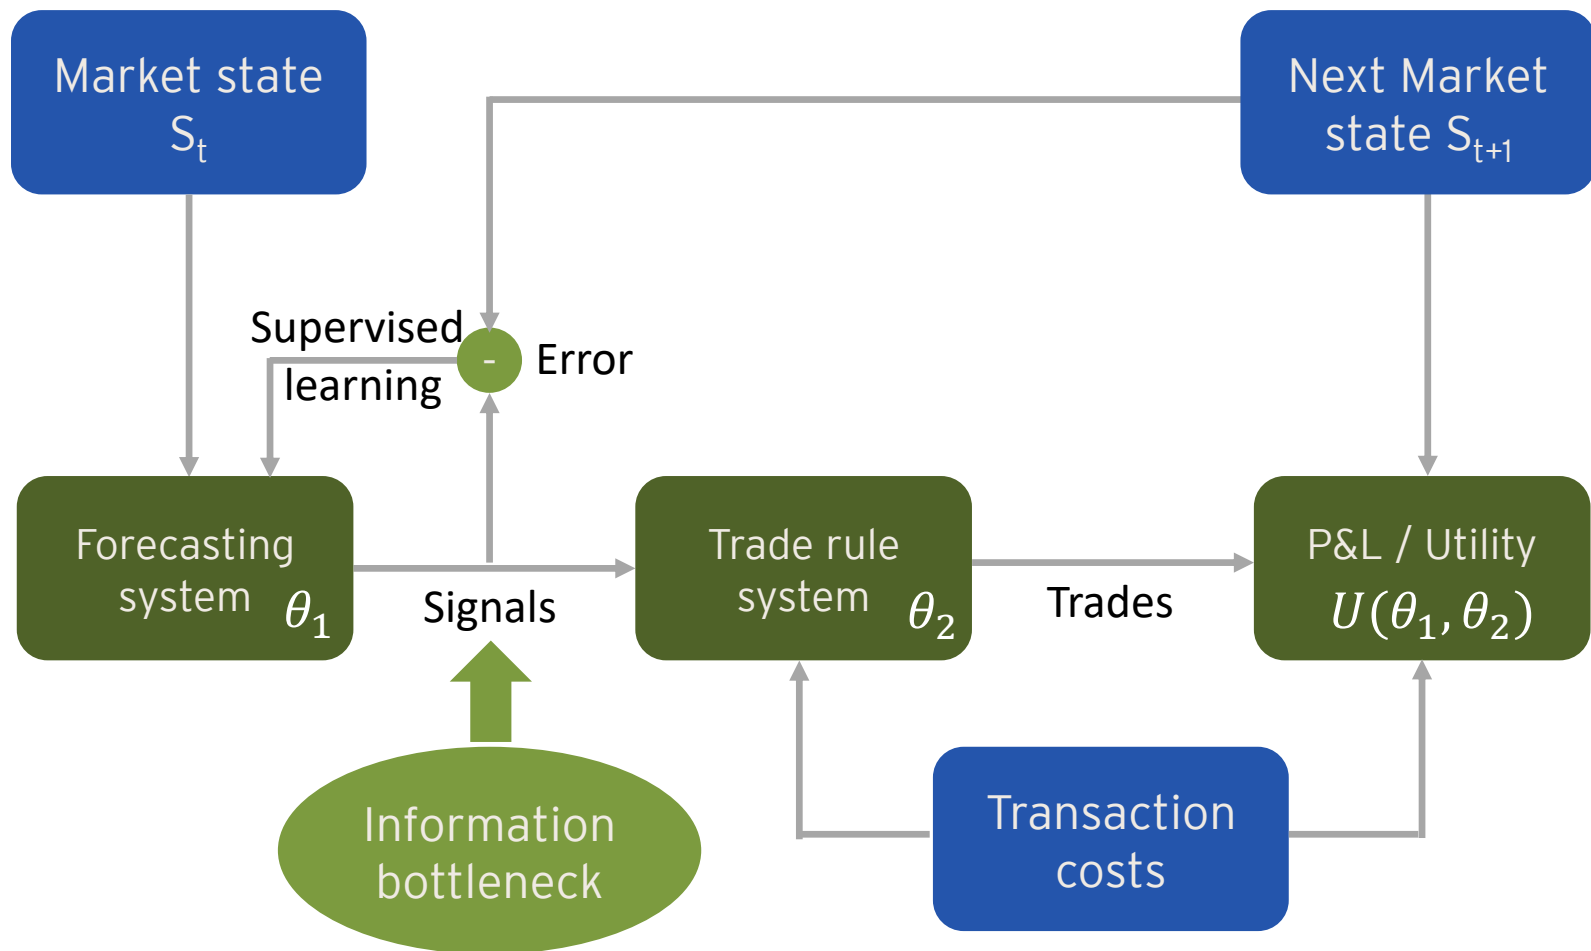
RL Algorithm Zoo

- Critic only
 - Q-learning
 - SARSA
- Actor only
 - Policy gradient
 - SRV
 - REINFORCE
- Actor-critic
 - TD(λ)
 - A2C, A3C
 - TRPO
 - PPO
 - DPPO

RL for Trading

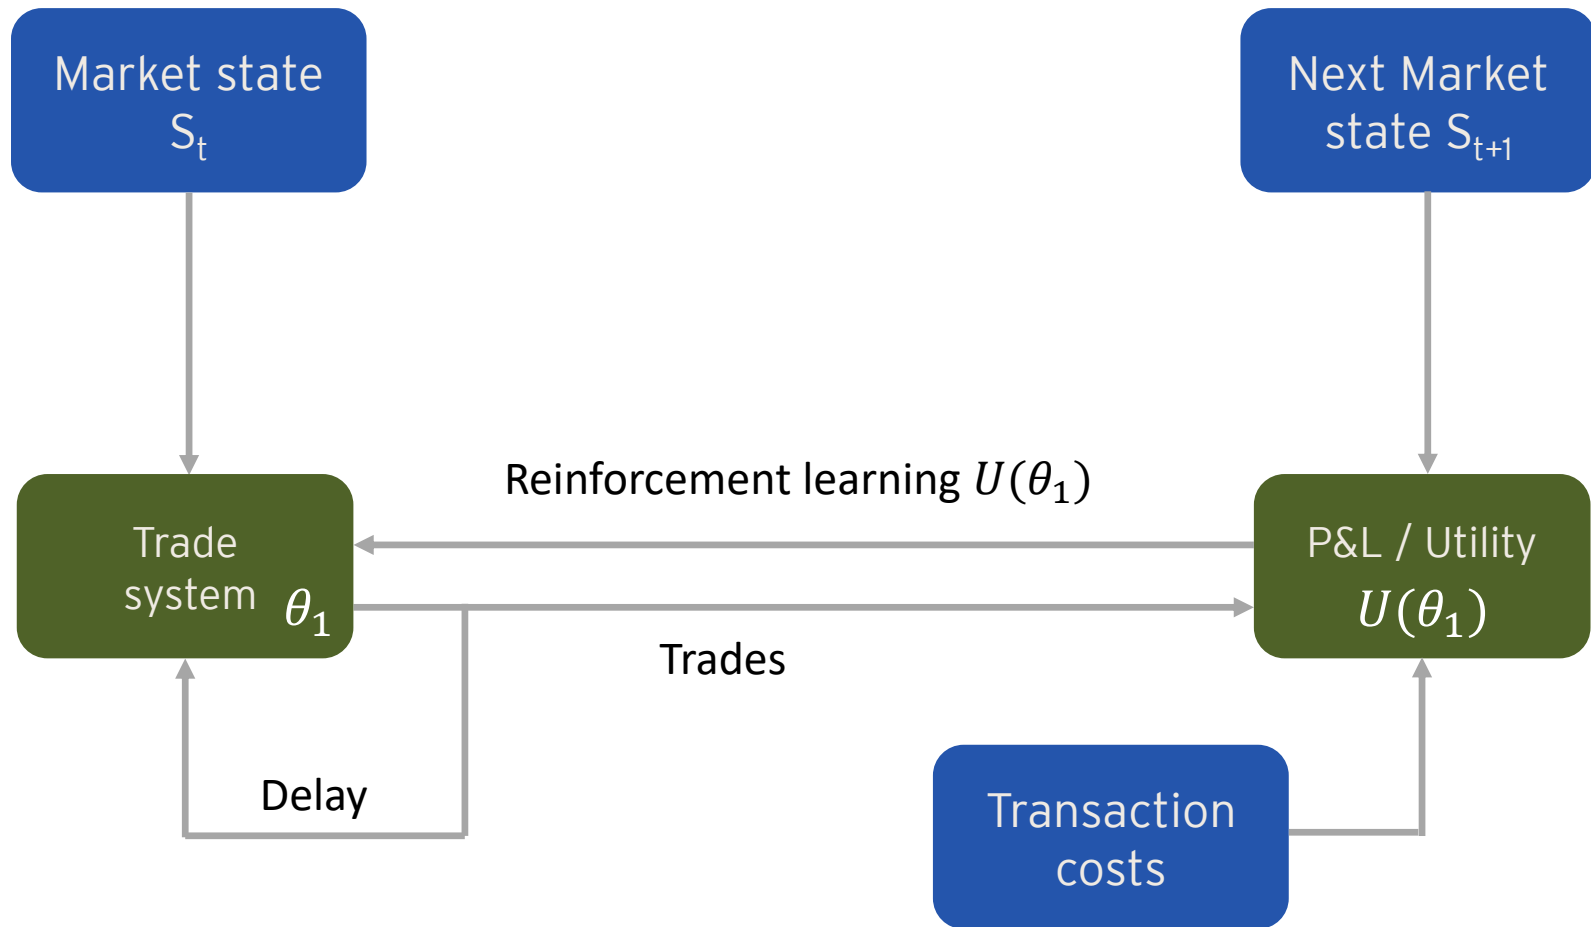
Why RL for Trading?

Signal based trading system



Why RL for Trading?

Reward driven trading system with reinforcement learning



Reinforcement learning based trading

- Single set of parameters
- Single utility function
- Utility includes transaction cost
- Direct mapping from market data to trades
- Easier to implement online updating
- Expensive model training

Signal based trading

- Parsimonious
- Two disconnected sets of parameters
- Supervised learning for signal forecast error is not utility
- Forecast system does not include transaction cost
- Two stage process causes information bottleneck

- How to model the state of the market and a trading system?
- How to handle time series data, tick data, order book data, ...?
- What additional time series data should be added?
- How to represent the policy? Discrete, continuous, hybrid...
- How to define the reward? Average, discounted
- How to include risk? Sharpe ratio,
- How to not forget catastrophic market events?
- How to incorporate additional constraints such as risk limits, ...?
- Which algorithms to use for model training?


Generative Models

Generative Models

- Why generative models?
 - What I cannot create, I do not understand (Richard Feynman)
- Two main classes
 - GANs – Generative adversarial networks
 - VAE – Variational autoencoders

Latent Variable Model

- Latent variable z can be thought of a encoded representation of x
- Likelihood serves as decoder
- Posterior provides encoder



Fit with Maximum Likelihood?

- Maximum likelihood standard model fitting approach

$$p(x) = \int p(x|z) p(z) dz \rightarrow \max$$

- **Problem:** marginal $p(x)$ and therefore also posterior


$$p(z|x) = \frac{p(x|z)p(z)}{p(x)}$$

are **intractable** and their calculation suffers from **exponential complexity**

- Solutions
 - Markov Chain MC, Hamiltonian MC
 - Approximation and variational inference

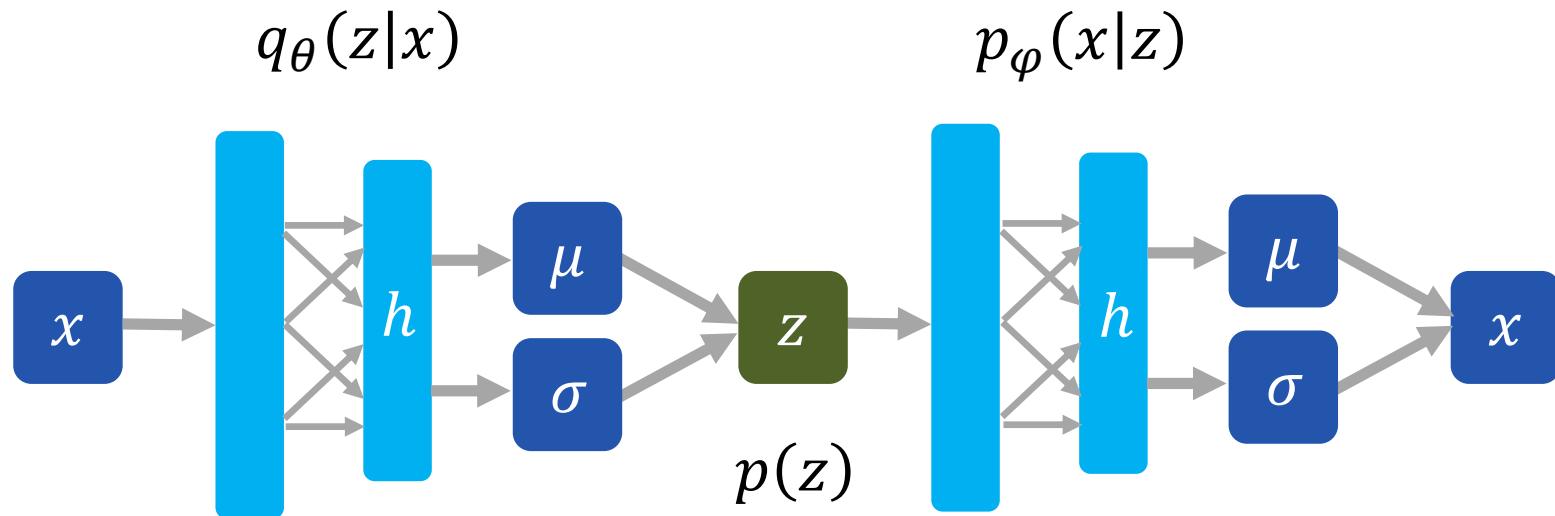
Variational Autoencoders

- Assume latent variable model with prior $p(z)$



Deep Variational Autoencoders

- Parameterize likelihood $p(x|z)$ with a deep neural network
- Approximate intractable posterior $p(z|x)$ with a deep neural network
- Learn the parameters θ and φ with backpropagation



Variational Inference

- Which loss to optimize?
- It turns out that there is a computable lower bound for the log of the marginal:

$$\underbrace{E_{q_{\theta}(z|x)}[\log p_{\varphi}(x, z)] - E_{q_{\theta}(z|x)}[\log q_{\theta}(z|x)]}_{ELBO(\theta, \varphi)} \leq \log p_{\varphi}(x)$$

$$ELBO(\theta, \varphi)$$

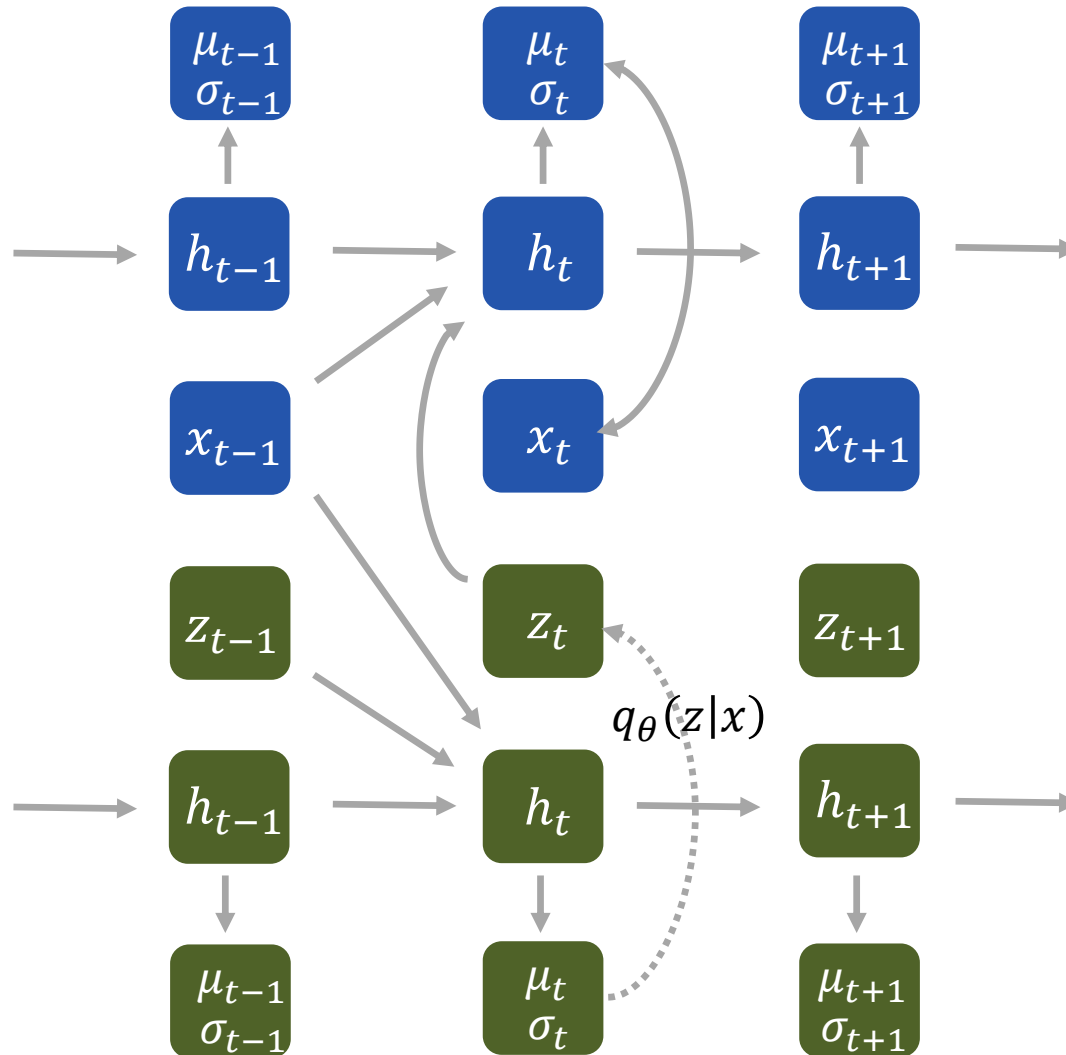
- Training criterion: maximize the evidence lower bound

$$ELBO(\theta, \varphi) \leq \log p_{\varphi}(x)$$

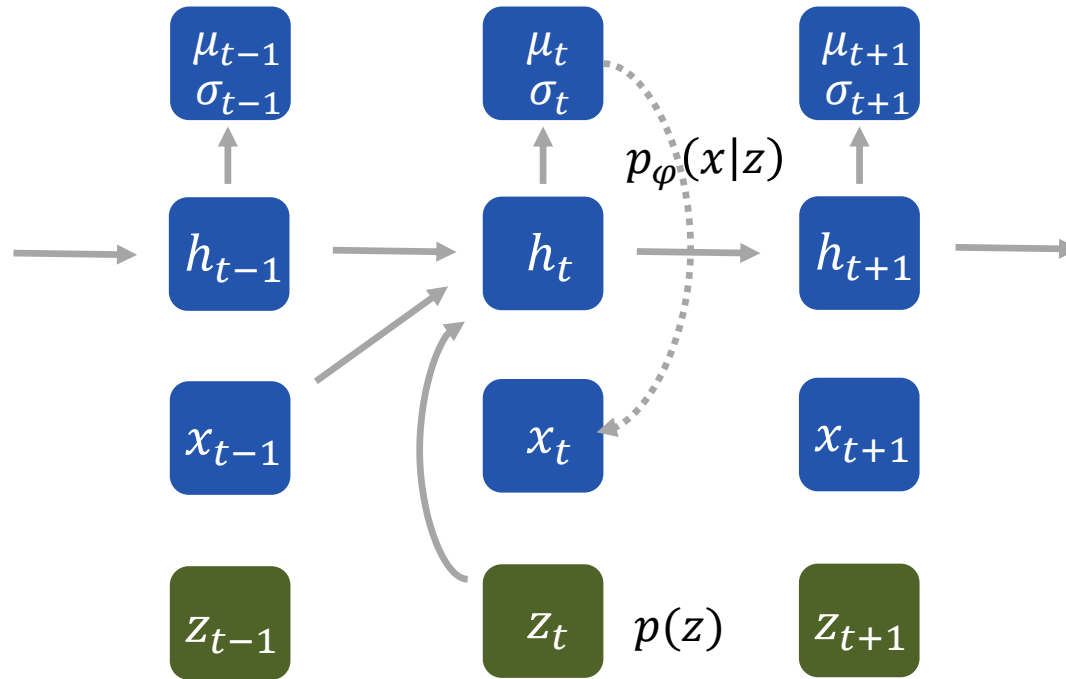
Applications to Time Series

- Sequence structure for observable and latent factor
- Model setup
 - Gaussian distributions with parameters calculated from deep recurrent neural network
 - Prior standard Gaussian
 - Model training with variational inference

Inference and Training


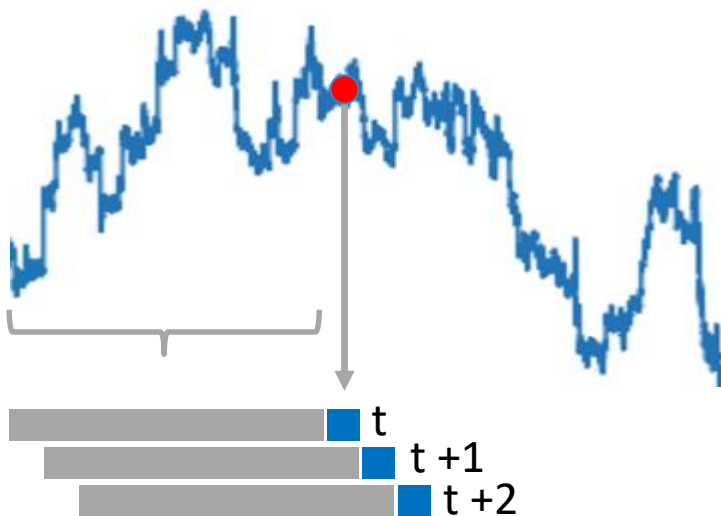


Generation



Time Series Embedding

- Single historical value not predictive enough
- Embedding
 - Use lag of 20 to 60 historical observations at every time step



TensorFlow Dynamic RNN

- Unrolling rnn with `tf.nn.dynamic_rnn`
 - Simple to use
 - Can handle variable sequence length
- Not flexible enough for generative networks

```
B = 3
D = 4
T = 5
PKEEP = 0.9

tf.reset_default_graph()

x = tf.placeholder(shape=[T, B, D], dtype=tf.float32)

with tf.variable_scope("RNN"):
    cell = tf.contrib.rnn.GRUCell(num_units = D)
    cell = tf.contrib.rnn.DropoutWrapper(cell, output_keep_prob = PKEEP)
    cells = tf.contrib.rnn.MultiRNNCell([cell])

    h, states = tf.nn.dynamic_rnn(cells, inputs = x, time_major=True, dtype=tf.float32)
```

- Using `tf.while_loop`
 - More to program, need to understand control structures in more detail
 - Much more flexible

```
x = tf.placeholder(shape=[T, B], dtype=tf.float32)

output_ta = tf.TensorArray(size=T, dtype=tf.float32)
input_ta = tf.TensorArray(size=T, dtype=tf.float32)
input_ta = input_ta.unstack(x)

def body(time, output_ta):
    xt = input_ta.read(time)
    output_ta = output_ta.write(time, tf.reduce_sum(xt**2))
    return (time+1, output_ta)

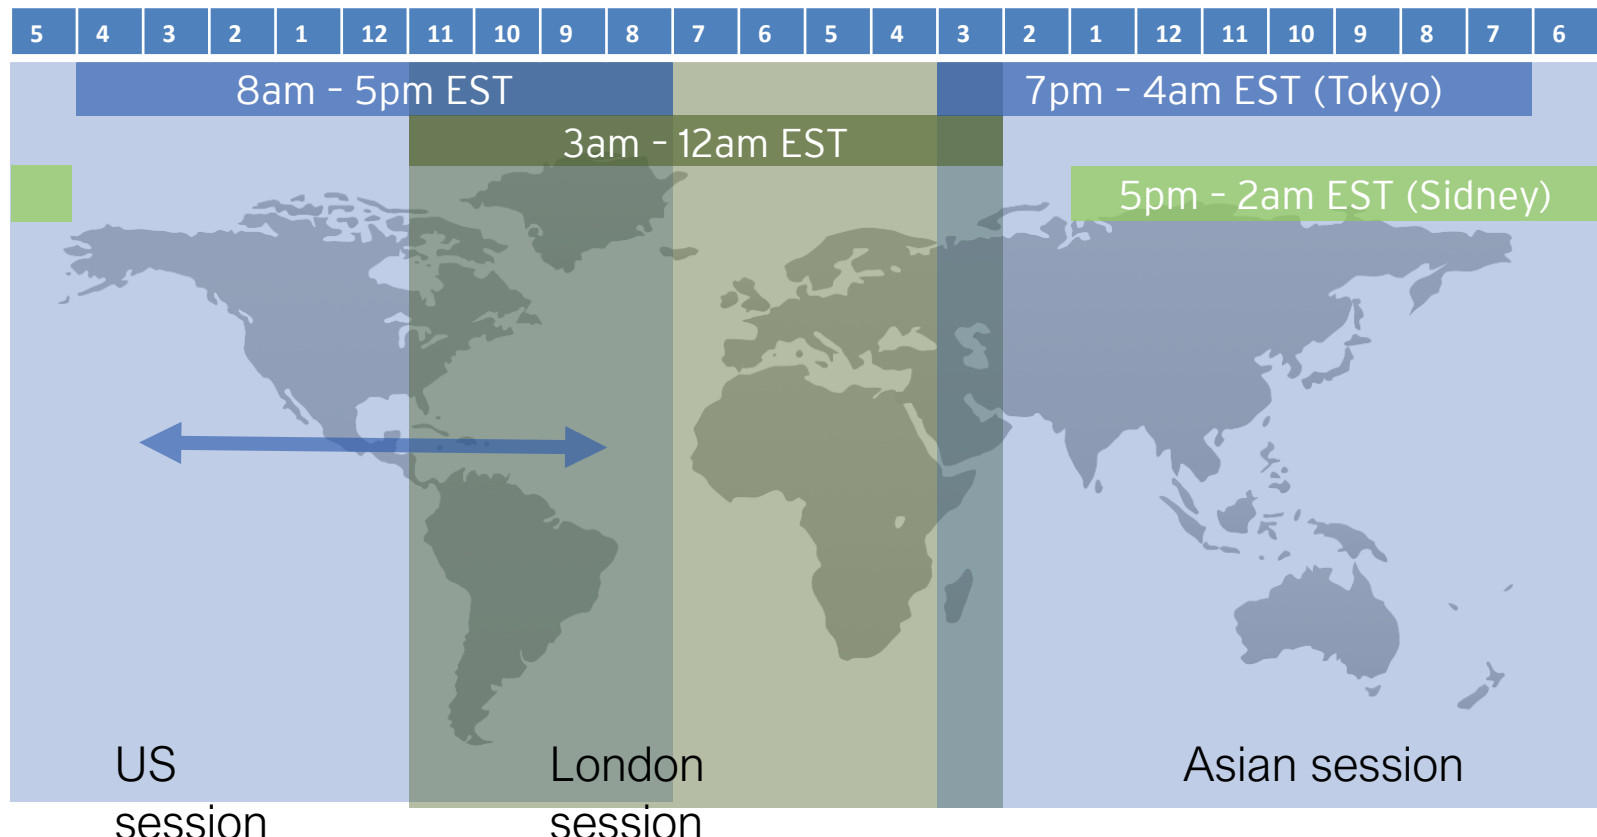
time_final, output_ta_final = tf.while_loop(
    cond=lambda time, *_ : time < T,
    body=body,
    loop_vars=(time, output_ta))

output_final = output_ta_final.stack()
```

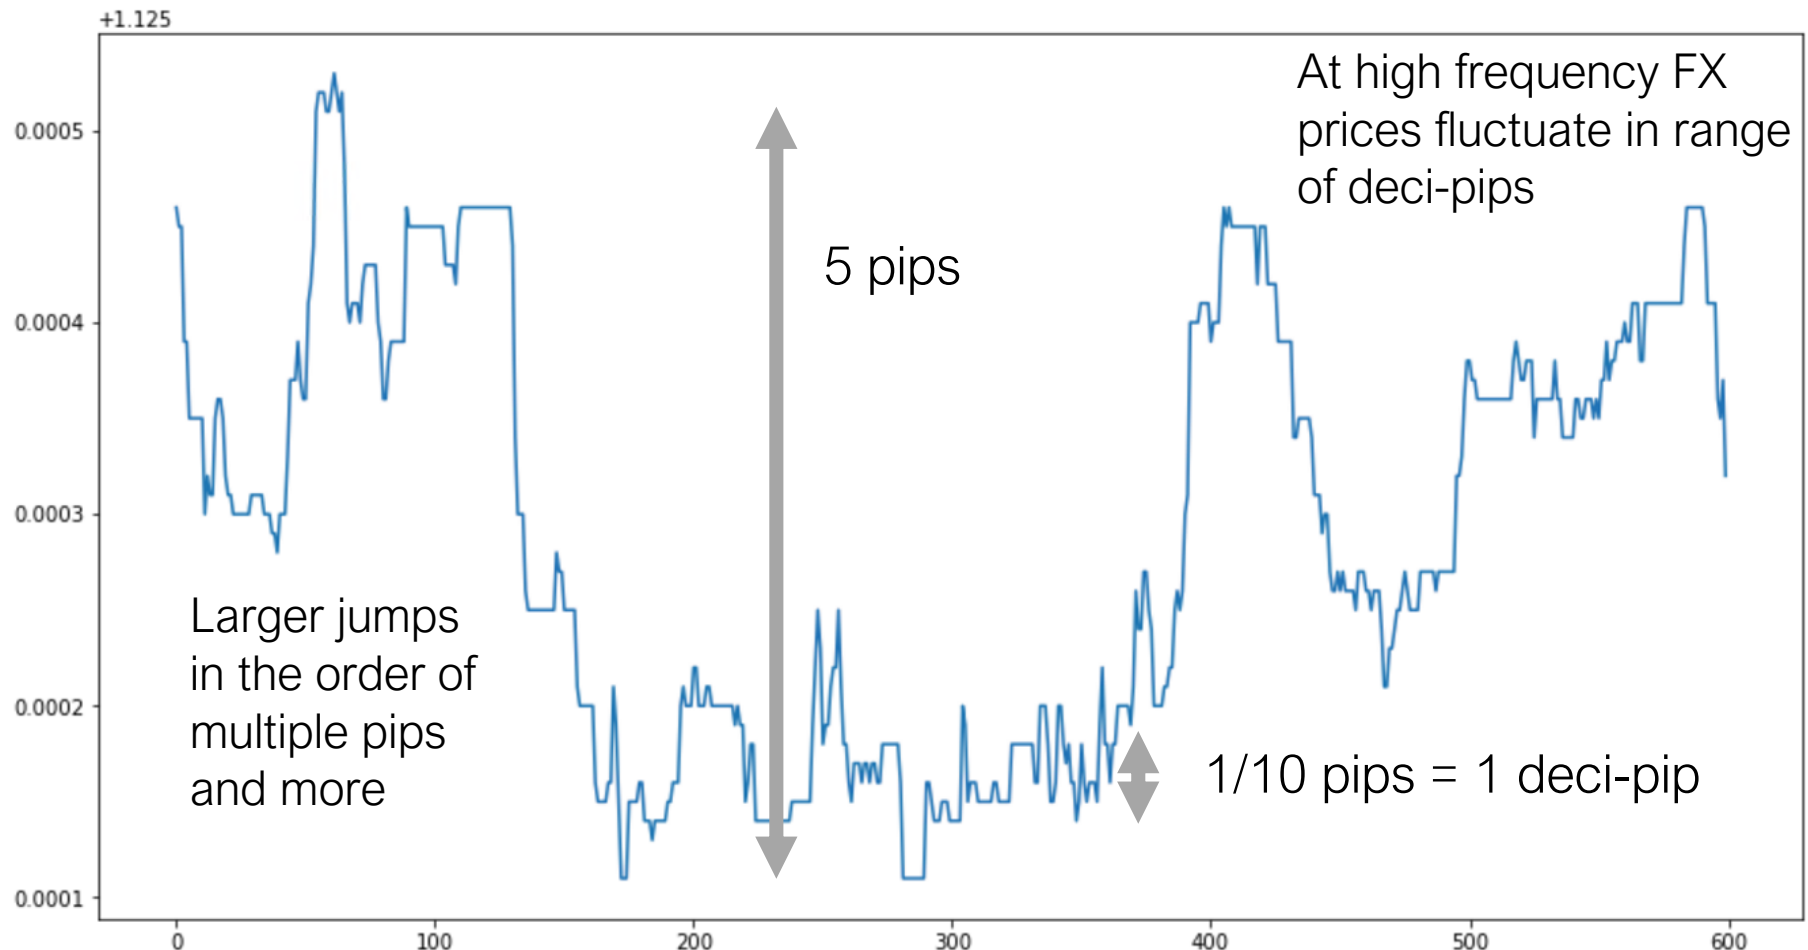
Application to FX Data

FX Data

- Collect tick data from major liquidity provider e.g. LMAX
- Aggregation to OHLC bars (1s, 10s, ...)
- Focus on US trading session

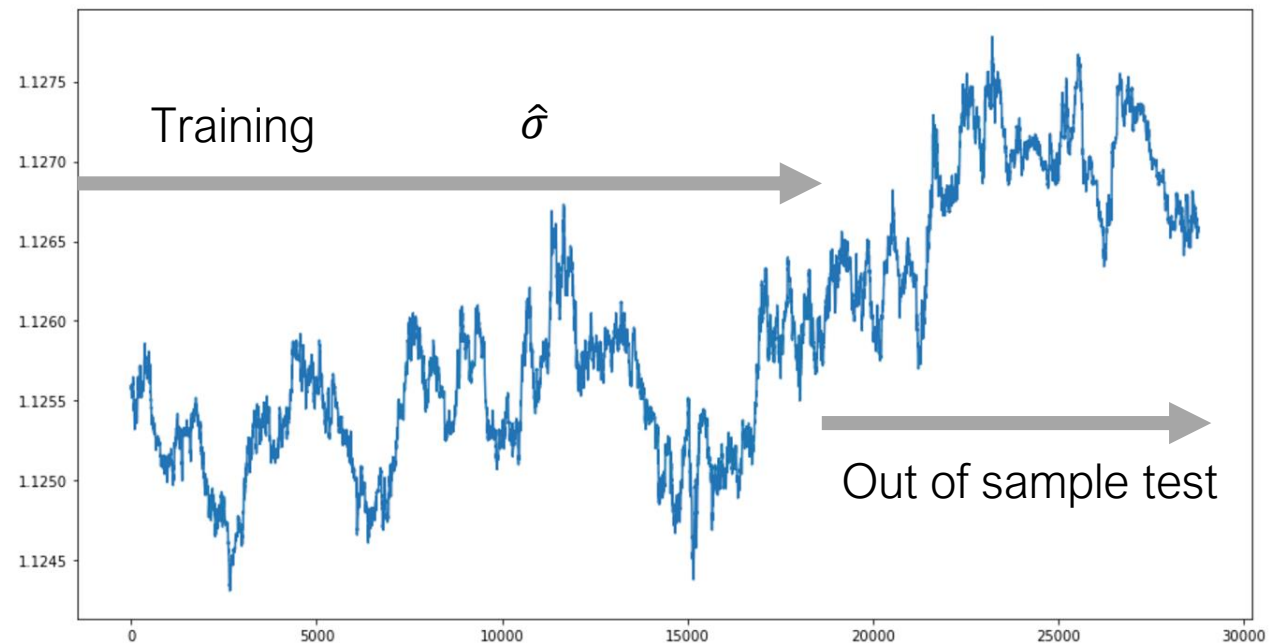


10 Min Sampled at 1s




Setup


- Normalize data with std deviation $\hat{\sigma}$ over training interval
- 260 trading days in 2016, one model per day
- 60 dim embedding, 2 dim latent space




Results



Out of Sample



Volatility of Prediction





@EgloffDaniel
@QuantAlea

Daniel Egloff
Dr. sc. math.
Phone: +41 79 430 03 61
daniel.egloff@quantalea.net